



R P Sarathy

Institute of Technology

An Autonomous Institution

Approved by AICTE | Accredited By NAAC A** | Affiliated to Anna University

Salem Bengaluru Highway NH - 7, Poosaripatty, Kadayampatty Taluk, Salem - 636305.

Admin Office: 93449-72274, Admission Cell: 93449-72275, 73977-56003,

Register No:

Name of the Laboratory:

Certified that this is bonafide record of work done by

_____ of _____ year/semester

_____ Branch for the year 20 - 20

Staff Incharge

Head of the Department

Submitted for the University Practical Examination

held on _____

Internal Examiner

External Examiner

INDEX

[illegible]

INDEX

[illegible]

EXPERIMENT: 1	Implementation of Uninformed search algorithms (BFS, DFS)
Date:	

Aim:

To implement and compare Uninformed Search Algorithms: Breadth-First Search (BFS) and Depth-First Search (DFS) for traversing or searching graph data structures.

1. Breadth-First Search (BFS)

Algorithm:

1. Start by adding the root node to a queue.
2. While the queue is not empty:
 - Dequeue a node from the front.
 - If the node is the goal, return the result.
 - Else, enqueue all its unvisited neighbors.
3. Mark each visited node to avoid processing it multiple times.

Program (Python):

```
1  from collections import deque
2
3  def bfs(graph, start, goal):
4      visited = set()
5      queue = deque([start])
6
7      while queue:
8          node = queue.popleft()
9
10         if node == goal:
11             return f"Goal {goal} found"
12
13         if node not in visited:
14             visited.add(node)
15             queue.extend(neighbor for neighbor in graph[node] if neighbor not in visited)
16
17     return f"Goal {goal} not found"
18
19 # Graph as an adjacency list
20 graph = {
21     'A': ['B', 'C'],
22     'B': ['D', 'E'],
23     'C': ['F'],
24     'D': [],
25     'E': ['F'],
26     'F': []
27 }
28
29 start_node = 'A'
30 goal_node = 'F'
31 print(bfs(graph, start_node, goal_node))
```

Output:

```
1  [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_1.py"
2  Goal F found
3
4  [Done] exited with code=0 in 0.083 seconds
5
```

2. Depth-First Search (DFS)

Algorithm:

1. Mark each visited node to avoid loops.
 - Else, push all its unvisited neighbors onto the stack.
 - If the node is the goal, return the result.
 - Pop a node from the stack.
2. While the stack is not empty:
3. Start by pushing the root node to a stack.

Program (Python):

```
1 def dfs(graph, start, goal):
2     visited = set()
3     stack = [start]
4
5     while stack:
6         node = stack.pop()
7
8         if node == goal:
9             return f"Goal {goal} found"
10
11        if node not in visited:
12            visited.add(node)
13            stack.extend(neighbor for neighbor in graph[node] if neighbor not in visited)
14
15        return f"Goal {goal} not found"
16
17 # Graph as an adjacency list (same as above)
18 print(dfs(graph, start_node, goal_node))
```

Output:

```
1 [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_1.py"
2 Goal F found
3
4 [Done] exited with code=0 in 0.083 seconds
```

Result:

EXPERIMENT: 2	Implementation of Uninformed search algorithms (BFS, DFS)
Date:	

Aim:

To implement and compare Informed Search Algorithms: A* and Memory-Bounded A* for finding the optimal path from a start node to a goal node in a graph.

1. A Search Algorithm*

Algorithm:

1. Initialize the open list (priority queue) with the start node.
2. Initialize the closed list as empty.
3. While the open list is not empty:
 - Select the node from the open list with the lowest $f(n)$ where $f(n) = g(n) + h(n)$, $g(n)$ is the cost from the start to n , and $h(n)$ is the heuristic estimate from n to the goal.
 - If the node is the goal, return the path.
 - Otherwise, expand the node by evaluating its neighbors, and add unvisited neighbors to the open list.
 - Move the current node to the closed list.

Program (Python):

```
1  import heapq
2
3  def a_star(graph, start, goal, h):
4      open_list = []
5      heapq.heappush(open_list, (0, start))
6      came_from = {}
7      g_score = {start: 0}
8      f_score = {start: h[start]}
9
10     while open_list:
11         _, current = heapq.heappop(open_list)
12
13         if current == goal:
14             return reconstruct_path(came_from, current)
15
16         for neighbor, cost in graph[current].items():
17             tentative_g = g_score[current] + cost
18             if neighbor not in g_score or tentative_g < g_score[neighbor]:
19                 came_from[neighbor] = current
20                 g_score[neighbor] = tentative_g
21                 f_score[neighbor] = tentative_g + h[neighbor]
22                 heapq.heappush(open_list, (f_score[neighbor], neighbor))
23
24     return "Path not found"
25
26 def reconstruct_path(came_from, current):
27     total_path = [current]
28     while current in came_from:
29         current = came_from[current]
30         total_path.append(current)
31     return total_path[::-1]
32
33 # Example graph and heuristic
34 graph = {
35     'A': {'B': 1, 'C': 4},
36     'B': {'D': 1, 'E': 3},
37     'C': {'E': 1},
38     'D': {'F': 2},
39     'E': {'F': 1},
40     'F': {}
41 }
42
43 heuristic = {
44     'A': 5, 'B': 4, 'C': 2,
45     'D': 2, 'E': 1, 'F': 0
46 }
47
48 start_node = 'A'
49 goal_node = 'F'
50 print(a_star(graph, start_node, goal_node, heuristic))
51
```

Output:

```
1 [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_2.py"
2 ['A', 'B', 'D', 'F']
3
4 [Done] exited with code=0 in 0.067 seconds
```

2. Memory-Bounded A (MA) Algorithm****Algorithm:**

1. Similar to A*, but limits the memory used by storing only a subset of nodes.
2. If memory exceeds a defined limit, remove the least promising nodes (with the highest $f(n)$ value) from the open list to free up space.
3. In case of pruning, store information about the best path to revisit later if necessary.

Program (Python):

```
1 import heapq
2
3 def memory_bounded_a_star(graph, start, goal, h, memory_limit):
4     open_list = []
5     heapq.heappush(open_list, (0, start))
6     came_from = {}
7     g_score = {start: 0}
8     f_score = {start: h[start]}
9     closed_list = set()
10
11     while open_list:
12         if len(open_list) > memory_limit:
13             # Prune the least promising node
14             open_list = sorted(open_list, key=lambda x: x[0])[:-1]
15
16         _, current = heapq.heappop(open_list)
17
18         if current == goal:
19             return reconstruct_path(came_from, current)
20
21         closed_list.add(current)
22
23         for neighbor, cost in graph[current].items():
24             tentative_g = g_score[current] + cost
25             if neighbor not in g_score or tentative_g < g_score[neighbor]:
26                 if neighbor not in closed_list:
27                     came_from[neighbor] = current
28                     g_score[neighbor] = tentative_g
29                     f_score[neighbor] = tentative_g + h[neighbor]
30                     heapq.heappush(open_list, (f_score[neighbor], neighbor))
31
32     return "Path not found"
33
34 # Example usage with memory limit
35 memory_limit = 4 # Limiting memory to store only 4 nodes at a time
36 print(memory_bounded_a_star(graph, start_node, goal_node, heuristic, memory_limit))
```

Output:

```
1 [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_2.py"
2 ['A', 'B', 'D', 'F']
3
4 [Done] exited with code=0 in 0.101 seconds
```

Result:

EXPERIMENT: 3	Implement naïve Bayes models
Date:	

Aim:

To implement a Naïve Bayes classifier that uses probability theory to predict the class of given data points based on input features, assuming independence among features.

Algorithm:

- 1. Calculate Prior Probabilities:**
Determine the probability of each class in the training data.
- 2. Compute Likelihood:**
For each feature and class, compute the likelihood using a probability distribution (e.g., Gaussian for continuous data).
- 3. Apply Bayes' Theorem:**
Combine the prior and likelihood to compute the posterior probability for each class.
- 4. Predict Class:**
Assign the data point to the class with the highest posterior probability.

Program (Python):

```
1  from sklearn.datasets import load_iris
2  from sklearn.model_selection import train_test_split
3  from sklearn.naive_bayes import GaussianNB
4  from sklearn.metrics import accuracy_score, classification_report
5
6  # Load dataset
7  iris = load_iris()
8  X, y = iris.data, iris.target
9
10 # Split data into training and test sets
11 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
12
13 # Initialize and train the Naïve Bayes model (GaussianNB for continuous data)
14 model = GaussianNB()
15 model.fit(X_train, y_train)
16
17 # Predict using the test set
18 y_pred = model.predict(X_test)
19
20 # Calculate accuracy and display classification report
21 accuracy = accuracy_score(y_test, y_pred)
22 report = classification_report(y_test, y_pred)
23
24 print("Accuracy:", accuracy)
25 print("Classification Report:\n", report)
26
```

Output:

```
1  [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_3.py"
2  Accuracy: 0.9777777777777777
3  Classification Report:
4
5      precision    recall  f1-score   support
6
7      0           1.00      1.00      1.00         19
8      1           1.00      0.92      0.96         13
9      2           0.93      1.00      0.96         13
10
11      accuracy          0.98
12      macro avg          0.98
13      weighted avg       0.98
14
15  [Done] exited with code=0 in 6.026 seconds
16
```

Result:

EXPERIMENT: 4	Implement Bayesian Networks
Date:	

Aim:

To implement a Bayesian Network that models probabilistic relationships among variables and performs inference to compute the likelihood of events given observed evidence.

Algorithm:

- 1. Define the Network Structure:**

Identify nodes (random variables) and their dependencies. For example, in the classic burglary scenario, nodes such as Burglary, Earthquake, Alarm, JohnCalls, and MaryCalls are connected by cause-effect relationships.

- 2. Specify Conditional Probability Distributions (CPDs):**

For each node, define the CPD. This includes prior probabilities for root nodes and conditional probabilities for nodes with parents.

- 3. Construct the Bayesian Network:**

Use a library (e.g., pgmpy) to create the network and add the nodes and CPDs.

- 4. Validate the Model:**

Check that the CPDs are consistent and that the network is correctly specified.

- 5. Perform Inference:**

Use an inference algorithm (e.g., Variable Elimination) to compute the posterior probability of a node given evidence from other nodes.

Program (Python):

```
1  # Import necessary libraries
2  from pgmpy.models import BayesianModel
3  from pgmpy.factors.discrete import TabularCPD
4  from pgmpy.inference import VariableElimination
5
6  # Define the structure of the Bayesian Network
7  model = BayesianModel([('Burglary', 'Alarm'),
8                        ('Earthquake', 'Alarm'),
9                        ('Alarm', 'JohnCalls'),
10                       ('Alarm', 'MaryCalls')])
11
12 # Define CPDs for each node
13
14 # Prior for Burglary (0: No burglary, 1: Burglary)
15 cpd_burglary = TabularCPD(variable='Burglary', variable_card=2,
16                           values=[[0.999], [0.001]])
17
18 # Prior for Earthquake (0: No earthquake, 1: Earthquake)
19 cpd_earthquake = TabularCPD(variable='Earthquake', variable_card=2,
20                              values=[[0.998], [0.002]])
21
22 # CPD for Alarm given Burglary and Earthquake
23 # The values are arranged for the combinations: [Burglary=0, Earthquake=0], [0,1], [1,0], [1,1]
24 cpd_alarm = TabularCPD(variable='Alarm', variable_card=2,
25                         values=[[0.999, 0.71, 0.06, 0.05], # Alarm = 0 (False)
26                                [0.001, 0.29, 0.94, 0.95]], # Alarm = 1 (True)
27                         evidence=['Burglary', 'Earthquake'],
28                         evidence_card=[2, 2])
29
30 # CPD for JohnCalls given Alarm
31 cpd_john = TabularCPD(variable='JohnCalls', variable_card=2,
32                       values=[[0.95, 0.10], # John does not call
33                              [0.05, 0.90]], # John calls
34                       evidence=['Alarm'], evidence_card=[2])
35
36 # CPD for MaryCalls given Alarm
37 cpd_mary = TabularCPD(variable='MaryCalls', variable_card=2,
38                       values=[[0.99, 0.30], # Mary does not call
39                              [0.01, 0.70]], # Mary calls
40                       evidence=['Alarm'], evidence_card=[2])
41
42 # Add CPDs to the model
43 model.add_cpds(cpd_burglary, cpd_earthquake, cpd_alarm, cpd_john, cpd_mary)
44
45 # Validate the model structure and CPDs
46 assert model.check_model(), "Model invalid: Check CPDs and structure."
47
48 # Perform inference using Variable Elimination
49 infer = VariableElimination(model)
50
51 # Example Query:
52 # Calculate the probability of a burglary given that both John and Mary called (1 indicates 'True')
53 query_result = infer.query(variables=['Burglary'], evidence={'JohnCalls': 1, 'MaryCalls': 1})
54 print(query_result)
55
```

Output:

```
1  +-----+-----+
2  | Burglary    | phi(Burglary) |
3  +=====+=====+
4  | Burglary(0) |          0.7158 |
5  +-----+-----+
6  | Burglary(1) |          0.2842 |
7  +-----+-----+
8
9  [Done] exited with code=0 in 13.645 seconds
```

Result:

EXPERIMENT: 5	Build Regression models
Date:	

Aim:

To build and evaluate regression models that predict continuous target values from input features. The example demonstrates a linear regression model to illustrate data preparation, model training, and evaluation.

Algorithm:

1. Data Preparation:

- Load and inspect the dataset.
- Preprocess the data (handle missing values, scale features if necessary).

2. Train-Test Split:

- Divide the dataset into training and testing subsets.

3. Model Training:

- Initialize a regression model (e.g., Linear Regression).
- Fit the model using the training data.

4. Model Evaluation:

- Predict target values on the test set.
- Compute evaluation metrics such as Mean Squared Error (MSE) and R-squared (coefficient of determination).

5. Result Analysis:

- Analyze the output metrics to determine the model's performance.

Program (Python):

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.datasets import fetch_california_housing
4 from sklearn.model_selection import train_test_split
5 from sklearn.linear_model import LinearRegression
6 from sklearn.metrics import mean_squared_error, r2_score
7
8 # Load the California Housing dataset
9 data = fetch_california_housing()
10 X = pd.DataFrame(data.data, columns=data.feature_names)
11 y = data.target
12
13 # Split the data into training and testing sets (70% train, 30% test)
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
15
16 # Initialize the Linear Regression model
17 model = LinearRegression()
18
19 # Train the model
20 model.fit(X_train, y_train)
21
22 # Predict the target for test set
23 y_pred = model.predict(X_test)
24
25 # Evaluate the model
26 mse = mean_squared_error(y_test, y_pred)
27 r2 = r2_score(y_test, y_pred)
28
29 # Output evaluation results
30 print("Mean Squared Error:", mse)
31 print("R-squared:", r2)
32
```

Output:

```
1 [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\tempCodeRunnerFile.py"
2 Mean Squared Error: 0.5305677824766755
3 R-squared: 0.5957702326061662
4
5 [Done] exited with code=0 in 16.576 seconds
```

Result:

EXPERIMENT: 6	Build decision trees and random forests
Date:	

Aim:

To build and compare Decision Tree and Random Forest models for classification tasks using a sample dataset, highlighting their performance and generalization capabilities.

Algorithm:

1. Data Preparation:

- Load a classification dataset (e.g., Iris).
- Split the dataset into training and testing sets.

2. Model Training:

- Train a Decision Tree classifier on the training data.
- Train a Random Forest classifier on the training data.

3. Model Evaluation:

- Predict class labels on the test set.
- Evaluate the models using metrics like accuracy and a classification report.

Program (Python):

```
1  from sklearn.datasets import load_iris
2  from sklearn.model_selection import train_test_split
3  from sklearn.tree import DecisionTreeClassifier
4  from sklearn.ensemble import RandomForestClassifier
5  from sklearn.metrics import accuracy_score, classification_report
6
7  # Load the Iris dataset
8  data = load_iris()
9  X, y = data.data, data.target
10
11 # Split the dataset (70% train, 30% test)
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
13
14 # Decision Tree Model
15 dt_model = DecisionTreeClassifier(random_state=42)
16 dt_model.fit(X_train, y_train)
17 dt_pred = dt_model.predict(X_test)
18
19 # Random Forest Model
20 rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
21 rf_model.fit(X_train, y_train)
22 rf_pred = rf_model.predict(X_test)
23
24 # Evaluation
25 print("Decision Tree Accuracy:", accuracy_score(y_test, dt_pred))
26 print("Decision Tree Classification Report:\n", classification_report(y_test, dt_pred))
27
28 print("Random Forest Accuracy:", accuracy_score(y_test, rf_pred))
29 print("Random Forest Classification Report:\n", classification_report(y_test, rf_pred))
30
```

Output:

```
1 [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\tempCodeRunnerFile.py"
2 Decision Tree Accuracy: 1.0
3 Decision Tree Classification Report:
4           precision    recall  f1-score   support
5
6      0           1.00      1.00      1.00        19
7      1           1.00      1.00      1.00        13
8      2           1.00      1.00      1.00        13
9
10     accuracy                1.00        45
11    macro avg           1.00      1.00      1.00        45
12    weighted avg           1.00      1.00      1.00        45
13
14 Random Forest Accuracy: 1.0
15 Random Forest Classification Report:
16           precision    recall  f1-score   support
17
18      0           1.00      1.00      1.00        19
19      1           1.00      1.00      1.00        13
20      2           1.00      1.00      1.00        13
21
22     accuracy                1.00        45
23    macro avg           1.00      1.00      1.00        45
24    weighted avg           1.00      1.00      1.00        45
25
26
27 [Done] exited with code=0 in 2.464 seconds
```

Result:

EXPERIMENT: 7	Build SVM models
Date:	

Aim:

To build and evaluate Support Vector Machine (SVM) models for classification tasks, demonstrating their effectiveness on datasets like the Iris dataset.

Algorithm:

1. Data Preparation:

- Load the dataset (e.g., Iris dataset).
- Split the data into training and testing sets.

2. Model Training:

- Initialize an SVM classifier with chosen hyperparameters.
- Train the SVM model on the training set.

3. Model Evaluation:

- Predict the labels for the test set.
- Evaluate the model using accuracy and classification metrics.

Program (Python):

```
1  from sklearn.datasets import load_iris
2  from sklearn.model_selection import train_test_split
3  from sklearn.svm import SVC
4  from sklearn.metrics import accuracy_score, classification_report
5
6  # Load the Iris dataset
7  data = load_iris()
8  X, y = data.data, data.target
9
10 # Split dataset into training and testing sets (70% train, 30% test)
11 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
12
13 # Initialize the SVM classifier with a linear kernel
14 svm_model = SVC(kernel='linear', random_state=42)
15
16 # Train the model
17 svm_model.fit(X_train, y_train)
18
19 # Predict using the test set
20 y_pred = svm_model.predict(X_test)
21
22 # Evaluate the model
23 accuracy = accuracy_score(y_test, y_pred)
24 report = classification_report(y_test, y_pred)
25
26 print("SVM Accuracy:", accuracy)
27 print("Classification Report:\n", report)
28
```

Output:

```
1  [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_7.py"
2  SVM Accuracy: 1.0
3  Classification Report:
4
5      precision    recall  f1-score   support
6
7      0           1.00      1.00      1.00         19
8      1           1.00      1.00      1.00         13
9      2           1.00      1.00      1.00         13
10
11      accuracy          1.00          1.00          1.00         45
12      macro avg          1.00          1.00          1.00         45
13      weighted avg          1.00          1.00          1.00         45
14
15  [Done] exited with code=0 in 1.607 seconds
```

Result:

EXPERIMENT: 8	Implement ensembling techniques
Date:	

Aim:

To implement ensemble techniques by combining multiple machine learning models (base learners) to improve prediction accuracy and robustness.

Algorithm:

1. Data Preparation:

- Load and preprocess the dataset.
- Split the data into training and testing sets.

2. Train Base Models:

- Initialize several base classifiers (e.g., Decision Tree, SVM, Logistic Regression).

3. Ensemble Model Construction:

- Combine base models using an ensemble strategy such as a Voting Classifier.

4. Model Evaluation:

- Train the ensemble model.
- Evaluate its performance on the test set using accuracy and other metrics.

Program (Python):

```
1  from sklearn.datasets import load_iris
2  from sklearn.model_selection import train_test_split
3  from sklearn.tree import DecisionTreeClassifier
4  from sklearn.svm import SVC
5  from sklearn.linear_model import LogisticRegression
6  from sklearn.ensemble import VotingClassifier
7  from sklearn.metrics import accuracy_score, classification_report
8
9  # Load the Iris dataset
10 data = load_iris()
11 X, y = data.data, data.target
12
13 # Split dataset into training and testing sets (70% train, 30% test)
14 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
15
16 # Initialize base models
17 dt_model = DecisionTreeClassifier(random_state=42)
18 svm_model = SVC(kernel='linear', probability=True, random_state=42)
19 lr_model = LogisticRegression(max_iter=200, random_state=42)
20
21 # Create the ensemble model using a Voting Classifier
22 ensemble_model = VotingClassifier(
23     estimators=[('dt', dt_model), ('svm', svm_model), ('lr', lr_model)],
24     voting='soft' # Use soft voting to average predicted probabilities
25 )
26
27 # Train the ensemble model
28 ensemble_model.fit(X_train, y_train)
29
30 # Predict using the test set
31 y_pred = ensemble_model.predict(X_test)
32
33 # Evaluate the model
34 accuracy = accuracy_score(y_test, y_pred)
35 report = classification_report(y_test, y_pred)
36
37 print("Ensemble Model Accuracy:", accuracy)
38 print("Classification Report:\n", report)
39
```

Output:

```
1 [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_8.py"
2 Ensemble Model Accuracy: 1.0
3 Classification Report:
4           precision    recall  f1-score   support
5
6      0           1.00      1.00      1.00        19
7      1           1.00      1.00      1.00        13
8      2           1.00      1.00      1.00        13
9
10     accuracy              1.00        45
11    macro avg              1.00      1.00      1.00        45
12 weighted avg              1.00      1.00      1.00        45
13
14
15 [Done] exited with code=0 in 1.665 seconds
```

Result:

EXPERIMENT: 9	Implement clustering algorithms
Date:	

Aim:

To implement clustering algorithms to group similar data points into clusters without prior knowledge of labels, demonstrating unsupervised learning using the K-Means clustering algorithm.

Algorithm:

1. Data Preparation:

- Load and preprocess the dataset (e.g., Iris dataset).

2. Clustering Process:

- **Initialization:** Randomly select k initial centroids.
- **Assignment:** Assign each data point to the nearest centroid.
- **Update:** Recalculate centroids as the mean of points in each cluster.
- **Iteration:** Repeat the assignment and update steps until convergence (i.e., centroids no longer change or a maximum iteration is reached).

3. Evaluation:

- Analyze the clustering output by reviewing cluster centers, inertia (sum of squared distances), and if applicable, visualizing clusters.

Program (Python):


```

1 import numpy as np
2 import pandas as pd
3 from sklearn.datasets import load_iris
4 from sklearn.cluster import KMeans
5 import matplotlib.pyplot as plt
6
7 # Load the Iris dataset
8 iris = load_iris()
9 X = iris.data
10
11 # Set the number of clusters (for Iris dataset, usually k=3 is appropriate)
12 k = 3
13
14 # Initialize and fit the KMeans clustering model
15 kmeans = KMeans(n_clusters=k, random_state=42)
16 kmeans.fit(X)
17
18 # Predict the cluster labels for the dataset
19 labels = kmeans.labels_
20 centroids = kmeans.cluster_centers_
21 inertia = kmeans.inertia_
22
23 # Print clustering results
24 print("Cluster Labels:", labels)
25 print("Cluster Centers:\n", centroids)
26 print("Inertia:", inertia)
27
28 # (Optional) Visualize clusters using the first two features
29 plt.scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis', marker='o', edgecolor='k')
30 plt.scatter(centroids[:, 0], centroids[:, 1], c='red', marker='x', s=200, label='Centroids')
31 plt.xlabel('Feature 1')
32 plt.ylabel('Feature 2')
33 plt.title('K-Means Clustering on Iris Dataset')
34 plt.legend()
35 plt.show()
36

```

Output:

[illegible]



Result:

EXPERIMENT: 10	Implement EM for Bayesian networks
Date:	

Aim:

To implement the Expectation-Maximization (EM) algorithm for parameter estimation in a Bayesian network with latent variables, using synthetic data to learn the conditional probability distributions.

Algorithm:

1. **Initialization:** Start with random guesses for parameters (e.g., prior probabilities and conditional probabilities).
2. **Expectation Step (E-Step):**
 - Compute the posterior distribution of latent variables given observed data and current parameters.
 - Calculate expected counts for latent variables using Bayes' theorem.
3. **Maximization Step (M-Step):**
 - Update parameters by maximizing the expected log-likelihood from the E-step.
 - Normalize expected counts to compute new probabilities.
4. **Convergence Check:** Repeat E and M steps until parameters stabilize (change < tolerance).

Program (Python):

```
1  import numpy as np
2
3  # Generate synthetic data
4  np.random.seed(42)
5  p_z_true = 0.6
6  p_x_given_z_true = {0: 0.3, 1: 0.8}
7  n_samples = 1000
8  Z = np.random.binomial(1, p_z_true, n_samples)
9  X = np.array([np.random.binomial(1, p_x_given_z_true[z]) for z in Z])
10
11 # Initialize parameters
12 p_z_current = 0.5
13 p_x_z0_current = 0.5
14 p_x_z1_current = 0.5
15 max_iter = 100
16 tolerance = 1e-4
17 log_likelihoods = []
18
19 for iteration in range(max_iter):
20     # E-Step: Compute posteriors  $P(Z=1 | X)$ 
21     e_z1 = []
22     for x in X:
23         if x == 1:
24             prob_z1 = p_z_current * p_x_z1_current
25             prob_z0 = (1 - p_z_current) * p_x_z0_current
26         else:
27             prob_z1 = p_z_current * (1 - p_x_z1_current)
28             prob_z0 = (1 - p_z_current) * (1 - p_x_z0_current)
29         total = prob_z0 + prob_z1
30         e_z1.append(prob_z1 / total if total != 0 else 0)
31     e_z1 = np.array(e_z1)
32
33     # M-Step: Update parameters
34     new_p_z = np.mean(e_z1)
35     numerator_x1_z0 = np.sum((X == 1) * (1 - e_z1))
36     denominator_z0 = np.sum(1 - e_z1)
37     new_p_x_z0 = numerator_x1_z0 / denominator_z0 if denominator_z0 != 0 else 0
38     numerator_x1_z1 = np.sum((X == 1) * e_z1)
39     denominator_z1 = np.sum(e_z1)
40     new_p_x_z1 = numerator_x1_z1 / denominator_z1 if denominator_z1 != 0 else 0
41
42     # Check convergence
43     deltas = [
44         abs(new_p_z - p_z_current),
45         abs(new_p_x_z0 - p_x_z0_current),
46         abs(new_p_x_z1 - p_x_z1_current)
47     ]
48     p_z_current, p_x_z0_current, p_x_z1_current = new_p_z, new_p_x_z0, new_p_x_z1
49
50     # Log likelihood
51     log_likelihood = 0
52     for x in X:
53         if x == 1:
54             term = (1 - p_z_current)*p_x_z0_current + p_z_current*p_x_z1_current
55         else:
56             term = (1 - p_z_current)*(1 - p_x_z0_current) + p_z_current*(1 - p_x_z1_current)
57         log_likelihood += np.log(term) if term != 0 else 0
58     log_likelihoods.append(log_likelihood)
59
60     if all(delta < tolerance for delta in deltas):
61         print(f"Converged at iteration {iteration + 1}")
62         break
63
64 print(f"Estimated P(Z=1): {p_z_current:.4f} (True: 0.6)")
65 print(f"Estimated P(X=1|Z=0): {p_x_z0_current:.4f} (True: 0.3)")
66 print(f"Estimated P(X=1|Z=1): {p_x_z1_current:.4f} (True: 0.8)")
```

Output:

An example output might be:

```
1  [Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_10.py"
2  Converged at iteration 2
3  Estimated P(Z=1): 0.5000 (True: 0.6)
4  Estimated P(X=1|Z=0): 0.6080 (True: 0.3)
5  Estimated P(X=1|Z=1): 0.6080 (True: 0.8)
6
7  [Done] exited with code=0 in 1.025 seconds
```

Result:

EXPERIMENT: 11	Build simple NN models
Date:	

Aim

To build a simple feedforward neural network (NN) model for classification using TensorFlow/Keras, trained on the MNIST dataset to recognize handwritten digits (0-9).

Algorithm

1. Data Preparation:

- Load MNIST dataset (28x28 grayscale images of digits).
- Normalize pixel values to [0, 1].
- Flatten images to 784-dimensional vectors.
- Split into training and test sets.

2. Model Architecture:

- Input layer: 784 neurons (one per pixel).
- Hidden layer: 128 neurons with ReLU activation.
- Output layer: 10 neurons (one per digit) with softmax activation.

3. Training:

- Loss: Sparse Categorical Crossentropy (for integer labels).
- Optimizer: Stochastic Gradient Descent (SGD).
- Metric: Accuracy.
- Train for 10 epochs.

4. Evaluation:

- Test accuracy calculation.

Program (Python):

```
1 import tensorflow as tf
2 from keras.models import Sequential
3 from keras.layers import Dense, Flatten
4 from keras.datasets import mnist
5 from keras.utils import to_categorical
6
7 # Load MNIST dataset
8 (x_train, y_train), (x_test, y_test) = mnist.load_data()
9
10 # Normalize the pixel values (scale between 0 and 1)
11 x_train, x_test = x_train / 255.0, x_test / 255.0
12
13 # One-hot encode the labels
14 y_train = to_categorical(y_train, num_classes=10)
15 y_test = to_categorical(y_test, num_classes=10)
16
17 # Build a simple Neural Network model
18 model = Sequential([
19     Flatten(input_shape=(28, 28)),
20     Dense(128, activation='relu'),
21     Dense(10, activation='softmax')
22 ])
23
24 # Compile the model with an optimizer and loss function
25 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
26
27 # Train the model with a validation split of 10%
28 model.fit(x_train, y_train, epochs=5, batch_size=32, validation_split=0.1)
29
30 # Evaluate the model on the test set
31 loss, accuracy = model.evaluate(x_test, y_test)
32 print("Test Loss:", loss)
33 print("Test Accuracy:", accuracy)
34
```

Output:

[Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_11.py"

Epoch 1/5

1/1688 [.....] - ETA: 9:53 - loss: 2.3978 - accuracy: 0.1250

34/1688 [.....] - ETA: 2s - loss: 1.4847 - accuracy: 0.6250

.....

1688/1688 [=====] - 4s 2ms/step - loss: 0.2789 -
accuracy: 0.9211 - val_loss: 0.1401 - val_accuracy: 0.9613

Epoch 2/5

1/1688 [.....] - ETA: 0s - loss: 0.1407 - accuracy: 0.9062
27/1688 [.....] - ETA: 3s - loss: 0.1097 - accuracy: 0.9641
.....
1688/1688 [=====] - 3s 2ms/step - loss: 0.1241 -
accuracy: 0.9634 - val_loss: 0.1061 - val_accuracy: 0.9677
Epoch 3/5

1/1688 [.....] - ETA: 2s - loss: 0.0572 - accuracy: 1.0000
32/1688 [.....] - ETA: 2s - loss: 0.0865 - accuracy: 0.9746
.....
1688/1688 [=====] - 3s 2ms/step - loss: 0.0849 -
accuracy: 0.9745 - val_loss: 0.0826 - val_accuracy: 0.9767
Epoch 4/5

1/1688 [.....] - ETA: 3s - loss: 0.1030 - accuracy: 0.9688
33/1688 [.....] - ETA: 2s - loss: 0.0796 - accuracy: 0.9773
.....
1688/1688 [=====] - 3s 2ms/step - loss: 0.0629 -
accuracy: 0.9807 - val_loss: 0.0853 - val_accuracy: 0.9747
Epoch 5/5

1/1688 [.....] - ETA: 2s - loss: 0.0758 - accuracy: 0.9688
35/1688 [.....] - ETA: 2s - loss: 0.0505 - accuracy: 0.9893
.....
313/313 [=====] - 0s 866us/step - loss: 0.0805 -
accuracy: 0.9748
Test Loss: 0.08050397038459778
Test Accuracy: 0.9747999906539917

[Done] exited with code=0 in 18.83 seconds

Result

EXPERIMENT: 12	Build deep learning NN models
Date:	

Aim

To build a deep learning neural network (DLNN) model using convolutional neural networks (CNNs) for image classification on the MNIST dataset, achieving higher accuracy than a simple feedforward NN.

Algorithm

1. Data Preparation:

- Load MNIST dataset (28x28 grayscale images).
- Reshape data to include channel dimension (required for CNNs).
- Normalize pixel values to [0, 1].
- One-hot encode labels.

2. Model Architecture:

- **Convolutional Layers:**
 - Conv2D (32 filters, 3x3 kernel, ReLU activation).
 - MaxPooling2D (2x2 pool size).
 - Conv2D (64 filters, 3x3 kernel, ReLU activation).
 - MaxPooling2D (2x2 pool size).
- **Dense Layers:**
 - Flatten layer to convert 2D features to 1D.
 - Dense layer (128 neurons, ReLU activation).
 - Dropout layer (0.5 rate for regularization).
 - Output layer (10 neurons, softmax activation).

3. Training:

- Loss: Categorical Crossentropy.
- Optimizer: Adam.
- Metrics: Accuracy.
- Train for 10 epochs with batch size 64.

4. Evaluation:

- Test accuracy and loss calculation.

Program (Python):

```
1 import tensorflow as tf
2 from keras import layers, models
3
4 # Load and preprocess data
5 mnist = tf.keras.datasets.mnist
6 (x_train, y_train), (x_test, y_test) = mnist.load_data()
7 x_train = x_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0
8 x_test = x_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0
9 y_train = tf.keras.utils.to_categorical(y_train, 10)
10 y_test = tf.keras.utils.to_categorical(y_test, 10)
11
12 # Build CNN model
13 model = models.Sequential([
14     layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
15     layers.MaxPooling2D((2, 2)),
16     layers.Conv2D(64, (3, 3), activation='relu'),
17     layers.MaxPooling2D((2, 2)),
18     layers.Flatten(),
19     layers.Dense(128, activation='relu'),
20     layers.Dropout(0.5),
21     layers.Dense(10, activation='softmax')
22 ])
23
24 # Compile and train
25 model.compile(optimizer='adam',
26               loss='categorical_crossentropy',
27               metrics=['accuracy'])
28 history = model.fit(x_train, y_train, epochs=10, batch_size=64,
29                     validation_data=(x_test, y_test))
30
31 # Evaluate
32 test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
33 print(f"Test accuracy: {test_acc * 100:.2f}%")
```

Output:

[Running] python -u "c:\Users\Admin\OneDrive\SEM 4\AIML\Code\EXE_12.py"

Epoch 1/10

1/938 [.....] - ETA: 7:48 - loss: 2.3010 - accuracy: 0.0625

.....

938/938 [=====] - 17s 18ms/step - loss: 0.2410 -
accuracy: 0.9268 - val_loss: 0.0495 - val_accuracy: 0.9838

Epoch 2/10

1/938 [.....] - ETA: 19s - loss: 0.0466 - accuracy: 0.9844

4/938 [.....] - ETA: 15s - loss: 0.0563 - accuracy: 0.9805

.....

938/938 [=====] - 18s 19ms/step - loss: 0.0862 -
accuracy: 0.9749 - val_loss: 0.0364 - val_accuracy: 0.9883

Epoch 3/10

1/938 [.....] - ETA: 29s - loss: 0.0491 - accuracy: 0.9688

4/938 [.....] - ETA: 16s - loss: 0.0711 - accuracy: 0.9766

.....

938/938 [=====] - 19s 20ms/step - loss: 0.0631 -
accuracy: 0.9813 - val_loss: 0.0306 - val_accuracy: 0.9889

Epoch 4/10

1/938 [.....] - ETA: 16s - loss: 0.0188 - accuracy: 1.0000

4/938 [.....] - ETA: 18s - loss: 0.0633 - accuracy: 0.9805

.....

938/938 [=====] - 19s 21ms/step - loss: 0.0498 -
accuracy: 0.9845 - val_loss: 0.0285 - val_accuracy: 0.9900

Epoch 5/10

1/938 [.....] - ETA: 20s - loss: 0.0220 - accuracy: 1.0000

4/938 [.....] - ETA: 16s - loss: 0.0394 - accuracy: 0.9883

.....

938/938 [=====] - 19s 20ms/step - loss: 0.0418 -
accuracy: 0.9875 - val_loss: 0.0247 - val_accuracy: 0.9915

Epoch 6/10

1/938 [.....] - ETA: 24s - loss: 0.0753 - accuracy: 0.9688
3/938 [.....] - ETA: 25s - loss: 0.0393 - accuracy: 0.9844
.....
937/938 [=====>.] - ETA: 0s - loss: 0.0356 - accuracy: 0.9893
938/938 [=====] - 20s 21ms/step - loss: 0.0356 - accuracy: 0.9893 - val_loss: 0.0215 - val_accuracy: 0.9917

Epoch 7/10

1/938 [.....] - ETA: 18s - loss: 0.0489 - accuracy: 0.9688
4/938 [.....] - ETA: 18s - loss: 0.0277 - accuracy: 0.9883
.....
938/938 [=====] - 20s 21ms/step - loss: 0.0315 - accuracy: 0.9903 - val_loss: 0.0225 - val_accuracy: 0.9923

Epoch 8/10

1/938 [.....] - ETA: 22s - loss: 0.0513 - accuracy: 0.9688
4/938 [.....] - ETA: 17s - loss: 0.0186 - accuracy: 0.9922
.....
938/938 [=====] - 21s 22ms/step - loss: 0.0277 - accuracy: 0.9913 - val_loss: 0.0233 - val_accuracy: 0.9923

Epoch 9/10

1/938 [.....] - ETA: 26s - loss: 0.0122 - accuracy: 1.0000
3/938 [.....] - ETA: 23s - loss: 0.0093 - accuracy: 1.0000
.....
938/938 [=====] - 19s 21ms/step - loss: 0.0221 - accuracy: 0.9929 - val_loss: 0.0255 - val_accuracy: 0.9922

Test accuracy: 99.22%

[Done] exited with code=0 in 196.858 seconds

Result